

Time Series Forecasting Using Attention-Transformers

Import Libraries

```
import math
import numpy as np
import pandas as pd

import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.layers import Input, Dense, Dropout, LayerNormalization, M

import yfinance as yf
```

```
def calculate_bollinger_bands(data, window=10, num_of_std=2):
    """Calculate Bollinger Bands"""
    rolling_mean = data.rolling(window=window).mean()
    rolling_std = data.rolling(window=window).std()
    upper_band = rolling_mean + (rolling_std * num_of_std)
    lower_band = rolling_mean - (rolling_std * num_of_std)
    return upper_band, lower_band

def calculate_rsi(data, window=10):
    """Calculate Relative Strength Index"""
    delta = data.diff()
    gain = delta.clip(lower=0)
    loss = -delta.clip(upper=0)
    avg_gain = gain.rolling(window=window, min_periods=1).mean()
    avg_loss = loss.rolling(window=window, min_periods=1).mean()
    rs = avg_gain / avg_loss
    rsi = 100 - (100 / (1 + rs))
    return rsi

def calculate_roc(data, periods=10):
    """Calculate Rate of Change."""
    roc = ((data - data.shift(periods)) / data.shift(periods)) * 100
    return roc
```

```
# List of tickers for the big tech companies, forming the MAMAA group
tickers = ['META', 'AAPL', 'MSFT', 'AMZN', 'GOOG']
```

```
ticker_data_frames = []
stats = {}
for ticker in tickers:

    # Download historical data for the ticker
    data = yf.download(ticker, period="60d", interval="5m")

    # Calculate the daily percentage change
    close = data['Close']
    upper, lower = calculate_bollinger_bands(close, window=14, num_of_std=2)
    width = upper - lower
    rsi = calculate_rsi(close, window=14)
    roc = calculate_roc(close, periods=14)
    volume = data['Volume']
    diff = data['Close'].diff(1)
    percent_change_close = data['Close'].pct_change() * 100

    # Create a DataFrame for the current ticker and append it to the list
    ticker_df = pd.DataFrame({
        ticker+'_close': close,
        ticker+'_width': width,
        ticker+'_rsi': rsi,
        ticker+'_roc': roc,
        ticker+'_volume': volume,
        ticker+'_diff': diff,
        ticker+'_percent_change_close': percent_change_close,
    })

    MEAN = ticker_df.mean()
    STD = ticker_df.std()

    # Keep track of mean and std
    for column in MEAN.index:
        stats[f"{column}_mean"] = MEAN[column]
        stats[f"{column}_std"] = STD[column]

    # Normalize the training features
    ticker_df = (ticker_df - MEAN) / STD

    ticker_data_frames.append(ticker_df)
```

```
# Convert the dictionary containing feature statistics to a DataFrame for easier
stats = pd.DataFrame([stats], index=[0])

# Display the DataFrame to verify its structure
stats.head()
```

```

META_close_mean META_close_std META_width_mean META_width_std META_rsi_mean META_rsi_std META_roc_mean META_roc_std META_volume_mean META_volume_std ...
0      487.84637      23.192329      5.729313      7.677995      50.634651      16.999532      -0.003481      1.103152      187859.55596      274559.281614 ...
1 rows x 70 columns

```

```
df = pd.concat(ticker_data_frames, axis=1)
df.replace([np.inf, -np.inf], np.nan, inplace=True)
df.dropna(inplace=True)
df.head()
```

```

DateTime META_close META_width META_rsi META_roc META_volume META_elff META_percent_change_close AAPL_close AAPL_width AAPL_rsi ... AMZN_volume AMZN_elff AMZN_percent_change_close GOOG_close GOOG_width GOOG_rsi GOOG_roc GOOG_volume GOOG_elff GOOG_percent_change_close
2024-02-08 10:40:00-05:00 -0.879676 -0.433255 0.270652 0.142999 -0.244458 0.602304 0.630348 2.153634 0.665285 -1.208763 ... 0.118234 0.641835 0.663281 -0.329635 -0.219043 -0.056453 -0.040230 -0.227587 0.076263 0.080303
2024-02-08 10:45:00-05:00 -0.878427 -0.466511 0.587228 0.258118 -0.071265 0.052056 0.052723 2.186818 0.630445 -1.293217 ... 0.080416 0.171885 0.178988 -0.307864 -0.220821 -0.077832 -0.058189 -0.127030 0.439402 0.456023
2024-02-08 10:50:00-05:00 -0.827704 -0.450117 1.183874 0.567584 0.383265 0.709802 0.740270 2.202234 0.427958 -1.099586 ... -0.178175 -0.134324 -0.138240 -0.286448 -0.190345 0.229790 0.097318 0.088416 0.194340 0.201217
2024-02-08 10:55:00-05:00 -0.800871 -0.354933 1.075216 0.485738 0.307327 0.390325 0.405340 2.207914 0.286212 -1.099265 ... 0.174715 0.167176 0.172114 -0.304069 -0.189304 0.205668 0.086381 0.074889 -0.128107 -0.133252
2024-02-08 11:00:00-05:00 -0.791053 -0.282378 1.239436 0.562948 0.085773 0.148033 0.180387 2.214404 0.161153 -0.808633 ... 0.073313 0.128162 0.129704 -0.307152 -0.215345 1.240011 0.442177 -0.048488 -0.089454 -0.093197
...
2024-05-03 15:35:00-04:00 -1.544859 -0.279655 0.799444 0.439825 0.047033 0.301050 0.324064 1.644218 0.334003 -1.173548 ... -0.058003 0.875828 0.830986 1.889887 -0.379528 0.863357 0.290154 -0.089033 0.858408 0.595442
2024-05-03 15:40:00-04:00 -1.534834 -0.309009 0.741985 0.398740 -0.038340 0.147801 0.187819 1.579919 0.362781 -1.550960 ... -0.058873 0.137479 0.128985 1.898638 -0.380749 0.724066 0.218817 -0.189971 0.207486 0.188985
2024-05-03 15:45:00-04:00 -1.557224 -0.386848 0.531897 0.309381 0.088132 -0.323338 -0.351159 1.538051 0.400186 -1.686002 ... -0.108711 -0.282724 -0.279124 1.890835 -0.382651 0.702740 0.213235 0.033718 -0.205491 -0.188070
2024-05-03 15:50:00-04:00 -1.567819 -0.429556 0.388045 0.235851 0.604889 -0.151978 -0.186178 1.457406 0.629032 -1.778367 ... 0.413864 0.493485 0.466863 1.826381 -0.320285 1.059723 0.417744 0.387008 0.877885 0.782297
2024-05-03 15:55:00-04:00 -1.538583 -0.459272 0.414547 0.252434 1.325258 0.431062 0.464633 1.441178 0.858829 -1.685598 ... 1.782702 -0.070221 -0.088045 1.911472 -0.296289 0.528884 0.187330 0.898633 -0.388041 -0.351019
4659 rows x 35 columns

```

```
# Shift the dataframe up by one to align current features with the next step's o
labels = df.shift(-1)

# Remove the last row from both the features and labels to maintain consistent d
df = df.iloc[:-1]
labels = labels.iloc[:-1]
```

```

SEQUENCE_LEN = 24 # 2 hours of data at 5-minute intervals

def create_sequences(data, labels, mean, std, sequence_length=SEQUENCE_LEN):
    sequences = []
    lab = []
    data_size = len(data)

    # Loop to create each sequence and its corresponding label
    for i in range(data_size - (sequence_length + 13)): # Ensure we have data for
        if i == 0:
            continue
        sequences.append(data[i:i + sequence_length]) # The sequence of data
        lab.append([labels[i-1], labels[i + 12], mean[0], std[0]]) # The label a

    return np.array(sequences), np.array(lab)

```

```

sequences_dict = {}
sequence_labels = {}
for ticker in tickers:

    # Extract close and volume data for the ticker
    close = df[ticker+'_close'].values
    width = df[ticker+'_width'].values
    rsi = df[ticker+'_rsi'].values
    roc = df[ticker+'_roc'].values
    volume = df[ticker+'_volume'].values
    diff = df[ticker+'_diff'].values
    pct_change = df[ticker+'_percent_change_close'].values

    # Combine close and volume data
    ticker_data = np.column_stack((close,
                                    width,
                                    rsi,
                                    roc,
                                    volume,
                                    diff,
                                    pct_change))

    # Generate sequences
    attribute = ticker+"_close"
    ticker_sequences, lab = create_sequences(ticker_data,
                                             labels[attribute].values[SEQUENCE_LEN:],
                                             stats[attribute+"_mean"].values,
                                             stats[attribute+"_std"].values)

    sequences_dict[ticker] = ticker_sequences
    sequence_labels[ticker] = lab

```

```
# Combine data and labels from all tickers
all_sequences = []
all_labels = []

for ticker in tickers:
    all_sequences.extend(sequences_dict[ticker])
    all_labels.extend(sequence_labels[ticker])

# Convert to numpy arrays
all_sequences = np.array(all_sequences)
all_labels = np.array(all_labels)
```

```
np.random.seed(42)
shuffled_indices = np.random.permutation(len(all_sequences))
all_sequences = all_sequences[shuffled_indices]
all_labels = all_labels[shuffled_indices]

train_size = int(len(all_sequences) * 0.9)

# Split sequences
train_sequences = all_sequences[:train_size]
train_labels = all_labels[:train_size]

other_sequences = all_sequences[train_size:]
other_labels = all_labels[train_size:]

shuffled_indices = np.random.permutation(len(other_sequences))
other_sequences = other_sequences[shuffled_indices]
other_labels = other_labels[shuffled_indices]

val_size = int(len(other_sequences) * 0.5)

validation_sequences = other_sequences[:val_size]
validation_labels = other_labels[:val_size]

test_sequences = other_sequences[val_size:]
test_labels = other_labels[val_size:]
```

```
def transformer_encoder(inputs, head_size, num_heads, ff_dim, dropout=0):
    # Attention and Normalization
    x = LayerNormalization(epsilon=1e-6)(inputs)
    x = MultiHeadAttention(key_dim=head_size, num_heads=num_heads, dropout=dropout)(x, x)
    x = Add()(x, inputs)

    # Feed Forward Part
    y = LayerNormalization(epsilon=1e-6)(x)
    y = Dense(ff_dim, activation="relu")(y)
    y = Dropout(dropout)(y)
    y = Dense(inputs.shape[-1])(y)
    return Add()([y, x])
```

```
def build_transformer_model(input_shape, head_size, num_heads, ff_dim, num_layers):
    inputs = Input(shape=input_shape)
    x = inputs
    for _ in range(num_layers):
        x = transformer_encoder(x, head_size, num_heads, ff_dim, dropout)
    x = GlobalAveragePooling1D()(x)
    x = LayerNormalization(epsilon=1e-6)(x)
    outputs = Dense(1, activation="linear")(x)
    return Model(inputs=inputs, outputs=outputs)
```

```
input_shape = train_sequences.shape[1:]
head_size = 256
num_heads = 16
ff_dim = 1024
num_layers = 12
dropout = 0.20

model = build_transformer_model(input_shape, head_size, num_heads, ff_dim, num_layers)
model.summary()
```

Total params: 1,708,558 (6.52 MB)
Trainable params: 1,708,558 (6.52 MB)
Non-trainable params: 0 (0.00 B)

```

def custom_mae_loss(y_true, y_pred):
    y_true_next = tf.cast(y_true[:, 1], tf.float64) # Extract the true next val
    y_pred_next = tf.cast(y_pred[:, 0], tf.float64) # Extract the predicted nex
    abs_error = tf.abs(y_true_next - y_pred_next) # Calculate the absolute erro
    return tf.reduce_mean(abs_error) # Return the mean of these errors

def dir_acc(y_true, y_pred):
    mean, std = tf.cast(y_true[:, 2], tf.float64), tf.cast(y_true[:, 3], tf.floa
    y_true_prev = (tf.cast(y_true[:, 0], tf.float64) * std) + mean # Un-scale p
    y_true_next = (tf.cast(y_true[:, 1], tf.float64) * std) + mean # Un-scale n
    y_pred_next = (tf.cast(y_pred[:, 0], tf.float64) * std) + mean # Un-scale p

    true_change = y_true_next - y_true_prev # Calculate true change
    pred_change = y_pred_next - y_true_prev # Calculate predicted change

    correct_direction = tf.equal(tf.sign(true_change), tf.sign(pred_change)) #
    return tf.reduce_mean(tf.cast(correct_direction, tf.float64)) # Return the

```



```

# Define a callback to save the best model
checkpoint_callback_train = ModelCheckpoint(
    "transformer_train_model.keras", # Filepath to save the best model
    monitor="dir_acc", # "loss", # Metric to monitor
    save_best_only=True, # Save only the best model
    mode="max", # Minimize the monitored metric
    verbose=1, # Display progress
)

# Define a callback to save the best model
checkpoint_callback_val = ModelCheckpoint(
    "transformer_val_model.keras", # Filepath to save the best model
    monitor="val_dir_acc", # "val_loss", # Metric to monitor
    save_best_only=True, # Save only the best model
    mode="max", # Minimize the monitored metric
    verbose=1, # Display progress
)

def get_lr_callback(batch_size=16, mode='cos', epochs=500, plot=False):
    lr_start, lr_max, lr_min = 0.0001, 0.005, 0.00001 # Adjust learning rate bc
    lr_ramp_ep = int(0.30 * epochs) # 30% of epochs for warm-up
    lr_sus_ep = max(0, int(0.10 * epochs) - lr_ramp_ep) # Optional sustain phase


    def lr_fn(epoch):
        if epoch < lr_ramp_ep: # Warm-up phase
            lr = (lr_max - lr_start) / lr_ramp_ep * epoch + lr_start
        elif epoch < lr_ramp_ep + lr_sus_ep: # Sustain phase at max learning rate
            lr = lr_max
        elif mode == 'cos':
            decay_total_epochs, decay_epoch_index = epochs - lr_ramp_ep - lr_sus_ep
            phase = math.pi * decay_epoch_index / decay_total_epochs
            lr = (lr_max - lr_min) * 0.5 * (1 + math.cos(phase)) + lr_min
        else:
            lr = lr_min # Default to minimum learning rate if mode is not recognized

    return lr

if plot: # Plot learning rate curve if plot is True
    plt.figure(figsize=(10, 5))
    plt.plot(np.arange(epochs), [lr_fn(epoch) for epoch in np.arange(epochs)])
    plt.xlabel('Epoch')
    plt.ylabel('Learning Rate')
    plt.title('Learning Rate Scheduler')
    plt.show()

return tf.keras.callbacks.LearningRateScheduler(lr_fn, verbose=True)

```

```
BATCH_SIZE = 64 # Number of training examples used to calculate each iteration'
EPOCHS = 100 # Total number of times the entire dataset is passed through the n

model.fit(
    train_sequences, # Training features
    train_labels, # Training labels
    validation_data=(validation_sequences, validation_labels), # Validation data
    epochs=EPOCHS, # Number of epochs to train for
    batch_size=BATCH_SIZE, # Size of each batch
    shuffle=True, # Shuffle training data before each epoch
    callbacks=[checkpoint_callback_train, checkpoint_callback_val, get_lr_callback_val]
)
```

```
Epoch 1: LearningRateScheduler setting learning rate to 0.0001.
Epoch 1/100
325/325 [=====] - ETA: 0s - loss: 0.2581 - dir_acc: 0.5365
Epoch 1: dir_acc improved from 0.53607 to 0.53647, saving model to transformer_train_model.keras

Epoch 1: val_dir_acc improved from 0.51855 to 0.52825, saving model to transformer_val_model.keras
325/325 [=====] - 50s 154ms/step - loss: 0.2581 - dir_acc: 0.5365 - val_loss: 0.2660 - val_dir_acc: 0.5283 - lr: 1.0000e-04

Epoch 2: LearningRateScheduler setting learning rate to 0.00026333333333333336.
Epoch 2/100
325/325 [=====] - ETA: 0s - loss: 0.2468 - dir_acc: 0.5404
Epoch 2: dir_acc improved from 0.53647 to 0.54041, saving model to transformer_train_model.keras

Epoch 2: val_dir_acc did not improve from 0.52825
325/325 [=====] - 50s 153ms/step - loss: 0.2468 - dir_acc: 0.5404 - val_loss: 0.2536 - val_dir_acc: 0.5146 - lr: 2.6333e-04

Epoch 3: LearningRateScheduler setting learning rate to 0.0004266666666666667.
Epoch 3/100
325/325 [=====] - ETA: 0s - loss: 0.2336 - dir_acc: 0.5329
Epoch 3: dir_acc did not improve from 0.54041

Epoch 3: val_dir_acc did not improve from 0.52825
325/325 [=====] - 49s 151ms/step - loss: 0.2336 - dir_acc: 0.5329 - val_loss: 0.2393 - val_dir_acc: 0.5237 - lr: 4.2667e-04

Epoch 4: LearningRateScheduler setting learning rate to 0.00059.
Epoch 4/100
325/325 [=====] - ETA: 0s - loss: 0.2266 - dir_acc: 0.5379
Epoch 4: dir_acc did not improve from 0.54041

Epoch 4: val_dir_acc improved from 0.52825 to 0.53482, saving model to transformer_val_model.keras
```

```

model.load_weights("transformer_val_model.keras") # Load the best model from th
accuracy = model.evaluate(test_sequences, test_labels)[1] # Evaluate the model
print(accuracy)

from sklearn.metrics import r2_score

predictions = model.predict(test_sequences) # Make predictions on the test data
r2 = r2_score(test_labels[:, 1], predictions[:, 0]) # Calculate R-squared value
print(f"R-squared: {r2}")

```

Fi

```

37/37 [=====] - 1s 39ms/step - loss: 0.1119 - dir_acc: 0.7044
0.7043918967247009
37/37 [=====] - 1s 33ms/step
R-squared: 0.9716192227979671

```

Final R-squared Loss : 0.97

